

# Movielens\_Project

November 23, 2023

```
[2]: #import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[3]: pip install matplotlib seaborn
```

Defaulting to user installation because normal site-packages is not writeable  
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/site-packages (3.6.3)  
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/site-packages (0.12.2)  
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/site-packages (from matplotlib) (1.0.7)  
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/site-packages (from matplotlib) (0.11.0)  
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/site-packages (from matplotlib) (4.33.3)  
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/site-packages (from matplotlib) (1.4.3)  
Requirement already satisfied: numpy>=1.19 in /usr/local/lib/python3.10/site-packages (from matplotlib) (1.23.5)  
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/site-packages (from matplotlib) (22.0)  
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/site-packages (from matplotlib) (9.1.1)  
Requirement already satisfied: pyparsing>=2.2.1 in /usr/local/lib/python3.10/site-packages (from matplotlib) (3.0.9)  
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/site-packages (from matplotlib) (2.8.2)  
Requirement already satisfied: pandas>=0.25 in /usr/local/lib/python3.10/site-packages (from seaborn) (1.5.3)  
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/site-packages (from pandas>=0.25->seaborn) (2022.1)  
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/site-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)

WARNING: There was an error checking the latest version of pip.

Note: you may need to restart the kernel to use updated packages.

```
[4]: # 1.Import the three datasets

import os

# Get the current working directory
current_directory = os.getcwd()

# Define the file names and extensions
movies_file_name = "movies.dat"
ratings_file_name = "ratings.dat"
users_file_name = "users.dat"

# Create the absolute file paths
movies_file = os.path.join(current_directory, movies_file_name)
ratings_file = os.path.join(current_directory, ratings_file_name)
users_file = os.path.join(current_directory, users_file_name)

# Define the column names for each dataset
movies_columns = ['MovieID', 'Title', 'Genres']
ratings_columns = ['UserID', 'MovieID', 'Rating', 'Timestamp']
users_columns = ['UserID', 'Gender', 'Age', 'Occupation', 'Zip-code']

# Read the datasets into pandas dataframes
movies_df = pd.read_csv(movies_file, sep='::', header=None,
    ↪names=movies_columns, encoding='latin-1', engine='python')
ratings_df = pd.read_csv(ratings_file, sep='::', header=None,
    ↪names=ratings_columns, encoding='latin-1', engine='python')
users_df = pd.read_csv(users_file, sep='::', header=None, names=users_columns,
    ↪encoding='latin-1', engine='python')

# Display the first few rows of each dataframe
print("Movies DataFrame:")
print(movies_df.head())

print("\nRatings DataFrame:")
print(ratings_df.head())

print("\nUsers DataFrame:")
print(users_df.head())
```

Movies DataFrame:

	MovieID	Title	Genres
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy

2	3	Grumpier Old Men (1995)	Comedy   Romance
3	4	Waiting to Exhale (1995)	Comedy   Drama
4	5	Father of the Bride Part II (1995)	Comedy

Ratings DataFrame:

	UserID	MovieID	Rating	Timestamp
0	1	1193	5	978300760
1	1	661	3	978302109
2	1	914	3	978301968
3	1	3408	4	978300275
4	1	2355	5	978824291

Users DataFrame:

	UserID	Gender	Age	Occupation	Zip-code
0	1	F	1	10	48067
1	2	M	56	16	70072
2	3	M	25	15	55117
3	4	M	45	7	02460
4	5	M	25	20	55455

```
[21]: movies_df.shape
```

```
[21]: (3883, 3)
```

```
[22]: ratings_df.shape
```

```
[22]: (1000209, 4)
```

```
[23]: users_df.shape
```

```
[23]: (6040, 5)
```

```
[ ]: # 2.Create a new dataset [Master_Data] with the following columns MovieID Title
      ↳ UserID Age Gender Occupation
      # Rating. (Hint: (i) Merge two tables at a time. (ii) Merge the tables using
      ↳ two primary keys MovieID & UserId)
```

```
[5]: # Get the current working directory
      current_directory = os.getcwd()

      # Merge ratings_df and users_df on 'UserID'
      merged_ratings_users = pd.merge(ratings_df, users_df, on='UserID')

      # Merge merged_ratings_users and movies_df on 'MovieID'
      master_data = pd.merge(merged_ratings_users, movies_df, on='MovieID')

      # Select the desired columns
```

```

master_data = master_data[['MovieID', 'Title', 'UserID', 'Age', 'Gender', 'Occupation', 'Rating']]

# Display the first few rows of the Master_Data dataframe
print("Master_Data DataFrame:")
print(master_data.head())

```

Master\_Data DataFrame:

	MovieID	Title	UserID	Age	Gender	\
0	1193	One Flew Over the Cuckoo's Nest (1975)	1	1	F	
1	1193	One Flew Over the Cuckoo's Nest (1975)	2	56	M	
2	1193	One Flew Over the Cuckoo's Nest (1975)	12	25	M	
3	1193	One Flew Over the Cuckoo's Nest (1975)	15	25	M	
4	1193	One Flew Over the Cuckoo's Nest (1975)	17	50	M	

	Occupation	Rating
0	10	5
1	16	5
2	12	4
3	7	4
4	1	5

```

[31]: # Print the column names of the DataFrame
print(master_data.columns)

```

```

Index(['MovieID', 'Title', 'UserID', 'Age', 'Gender', 'Occupation', 'Rating'],
      dtype='object')

```

```

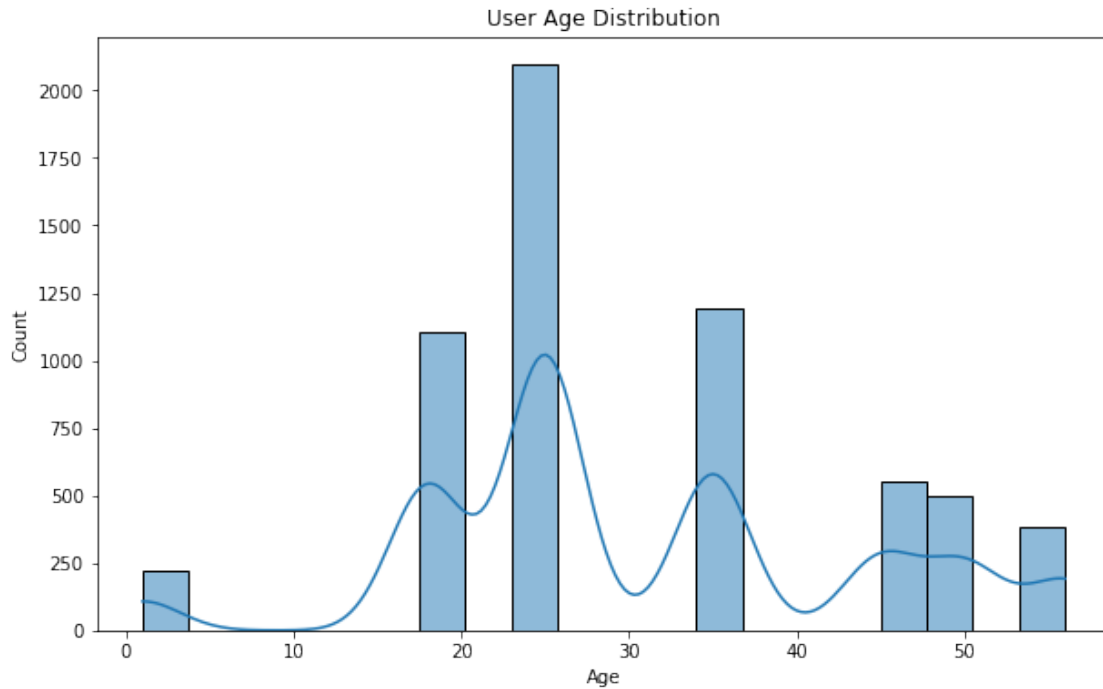
[ ]: # 3.Explore the datasets using visual representations (graphs or tables),
      also include your comments on the following:

```

```

[28]: # 3.1 User Age Distribution
# Plot the User Age Distribution
plt.figure(figsize=(10, 6))
sns.histplot(users_df['Age'], bins=20, kde=True)
plt.title('User Age Distribution')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()

```




```
[23]: # 3.2 User rating of the movie "Toy Story"

toystoryRating = master_data[master_data['Title'].str.contains('Toy Story')]

# Group by Title and Rating, count occurrences, and plot
toystoryRating.groupby(["Title", "Rating"]).size().unstack().plot(
    kind='barh', stacked=False, legend=True
)

plt.title('User Rating of the Movie "Toy Story"')
plt.xlabel('Count')
plt.ylabel('Rating')
plt.show()
```




```
[25]: # 3.3 Top 25 movies by viewership rating
# Group by MovieID and calculate the mean rating for each movie
movie_ratings = master_data.groupby('MovieID')['Rating'].mean()

# Sort movies by mean rating in descending order
top_movies = movie_ratings.sort_values(ascending=False).head(25)

# Get the movie titles for the top 25 movies
top_movies_titles = master_data[master_data['MovieID'].isin(top_movies.
    ↪index)]['Title'].unique()

# Plot the top 25 movies by viewership rating
plt.figure(figsize=(12, 8))
sns.barplot(x=top_movies, y=top_movies_titles, palette='viridis')
plt.title('Top 25 Movies by Viewership Rating')
plt.xlabel('Average Rating')
plt.ylabel('Movie Title')
plt.show()
```



[27]: # 3.4 Find the ratings for all the movies reviewed by for a particular user of `user id = 2696`

```
# Filter the master_data DataFrame for user ID 2696
user_2696_ratings = master_data[master_data['UserID'] == 2696][['Title',
    'Rating']]

# Display the ratings for the movies reviewed by user 2696 as a table
print(user_2696_ratings.to_markdown(index=False))
```

Title	Rating
Back to the Future (1985)	2
E.T. the Extra-Terrestrial (1982)	3
L.A. Confidential (1997)	4
Lone Star (1996)	5
JFK (1991)	1
Talented Mr. Ripley, The (1999)	4
Midnight in the Garden of Good and Evil (1997)	4
Cop Land (1997)	3
Palmetto (1998)	4
Perfect Murder, A (1998)	4
Game, The (1997)	4
I Know What You Did Last Summer (1997)	2
Devil's Advocate, The (1997)	4
Psycho (1998)	4
Wild Things (1998)	4
Basic Instinct (1992)	4
Lake Placid (1999)	1
Shining, The (1980)	4

I Still Know What You Did Last Summer (1998)		2	
Client, The (1994)		3	

[ ]: # 4. Feature Engineering:

```
[36]: # 4.1 Find out all the unique genres (Hint: split the data in column genre,
      ↪making a list and
      # then process the data to find out only the unique categories of genres)

      # Split the genres column into lists of genres
      genres_lists = movies_df['Genres'].str.split('|')

      # Create a set to store unique genres
      unique_genres = set()

      # Iterate over the lists of genres and add unique genres to the set
      for genres_list in genres_lists:
          unique_genres.update(genres_list)

      # Convert the set of unique genres to a list
      unique_genres_list = list(unique_genres)

      # Display the unique genres
      print("Unique Genres:")
      for genre in unique_genres_list:
          print(genre)
```

Unique Genres:

Crime  
 Horror  
 Animation  
 Children's  
 Adventure  
 Musical  
 Thriller  
 Documentary  
 War  
 Mystery  
 Comedy  
 Western  
 Fantasy  
 Sci-Fi  
 Romance  
 Film-Noir  
 Action  
 Drama



```
[34]: # Split the genres column into lists of genres
genres_lists = movies_df['Genres'].str.split('|')

# Create a set to store unique genres
unique_genres = set()

# Iterate over the lists of genres and add unique genres to the set
for genres_list in genres_lists:
    unique_genres.update(genres_list)

# Convert the set of unique genres to a list
unique_genres_list = list(unique_genres)

# Display the unique genres
print("Unique Genres:")
for genre in unique_genres_list:
    print(genre)
```

Unique Genres:

Crime  
Horror  
Animation  
Children's  
Adventure  
Musical  
Thriller  
Documentary  
War  
Mystery  
Comedy  
Western  
Fantasy  
Sci-Fi  
Romance  
Film-Noir  
Action  
Drama

```
[38]: # 4.2 Create a separate column for each genre category with a one-hot encoding
      ↪ ( 1 and 0)
# whether or not the movie belongs to that genre.# If 'master_data' does not
      ↪ have 'Genres', merge it with 'movies_df'
if 'Genres' not in master_data.columns:
    master_data = pd.merge(master_data, movies_df[['MovieID', 'Genres']],
      ↪ on='MovieID', how='left')

# Assuming master_data has a 'Genres' column
```

```

genres_lists = master_data['Genres'].str.split('|')

# One-hot encode genres
one_hot_encoded_genres = pd.get_dummies(genres_lists.apply(pd.Series).stack()).
    groupby(level=0).sum()

# Concatenate one-hot encoded genres with master_data
master_data = pd.concat([master_data, one_hot_encoded_genres], axis=1)

# Display the updated master_data
print(master_data)

```

	MovieID	Title	UserID	Age	\
0	1193	One Flew Over the Cuckoo's Nest (1975)	1	1	
1	1193	One Flew Over the Cuckoo's Nest (1975)	2	56	
2	1193	One Flew Over the Cuckoo's Nest (1975)	12	25	
3	1193	One Flew Over the Cuckoo's Nest (1975)	15	25	
4	1193	One Flew Over the Cuckoo's Nest (1975)	17	50	
...	...	...	...	...	
1000204	2198	Modulations (1998)	5949	18	
1000205	2703	Broken Vessels (1998)	5675	35	
1000206	2845	White Boys (1999)	5780	18	
1000207	3607	One Little Indian (1973)	5851	18	
1000208	2909	Five Wives, Three Secretaries and Me (1998)	5938	25	

	Gender	Occupation	Rating	Genres	Action	Adventure	\
0	F	10	5	Drama	0	0	
1	M	16	5	Drama	0	0	
2	M	12	4	Drama	0	0	
3	M	7	4	Drama	0	0	
4	M	1	5	Drama	0	0	
...	...	...	...	...	...	...	
1000204	M	17	5	Documentary	0	0	
1000205	M	14	3	Drama	0	0	
1000206	M	17	1	Drama	0	0	
1000207	F	20	5	Comedy Drama Western	0	0	
1000208	M	1	4	Documentary	0	0	

	...	Fantasy	Film-Noir	Horror	Musical	Mystery	Romance	Sci-Fi	\
0	...	0	0	0	0	0	0	0	
1	...	0	0	0	0	0	0	0	
2	...	0	0	0	0	0	0	0	
3	...	0	0	0	0	0	0	0	
4	...	0	0	0	0	0	0	0	
...	...	...	...	...	...	...	...	...	
1000204	...	0	0	0	0	0	0	0	
1000205	...	0	0	0	0	0	0	0	

1000206	...	0	0	0	0	0	0	0
1000207	...	0	0	0	0	0	0	0
1000208	...	0	0	0	0	0	0	0

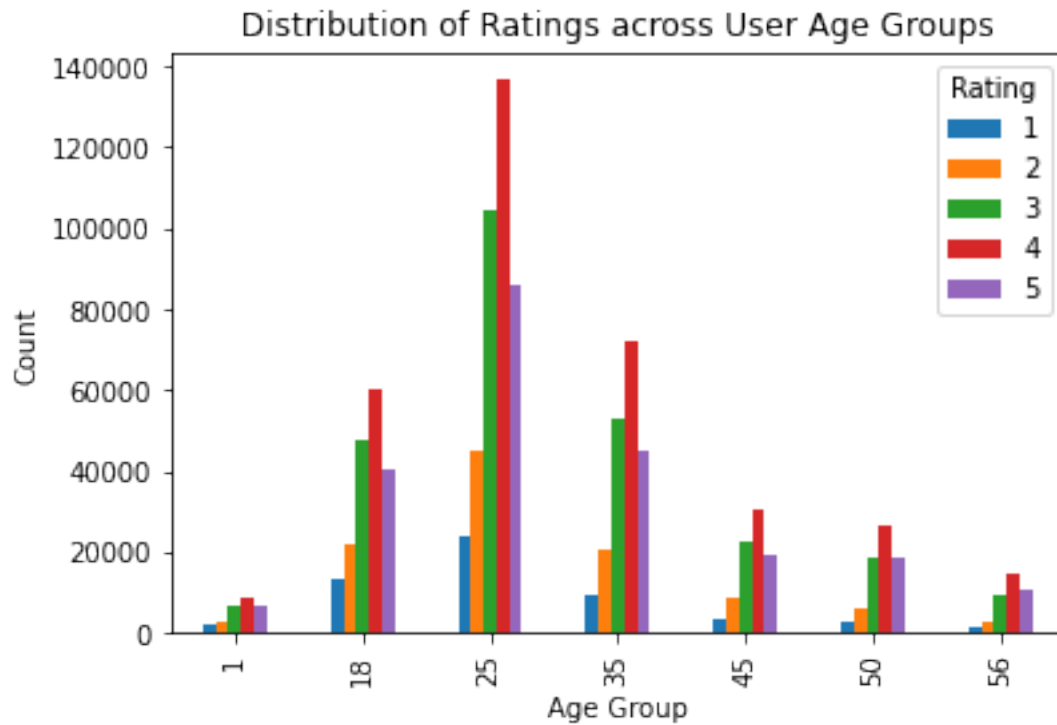
		Thriller	War	Western
0		0	0	0
1		0	0	0
2		0	0	0
3		0	0	0
4		0	0	0
...	...	...	...	...
1000204		0	0	0
1000205		0	0	0
1000206		0	0	0
1000207		0	0	1
1000208		0	0	0

[1000209 rows x 44 columns]

[42]: *# 4.3 Determine the features affecting the ratings of any particular movie.*


```
# Explore the relationship between user age and ratings
plt.figure(figsize=(12, 6))
ratings_by_age = master_data.groupby(['Age', 'Rating']).size().unstack()
ratings_by_age.plot(kind='bar', stacked=False, legend=True)
plt.title('Distribution of Ratings across User Age Groups')
plt.xlabel('Age Group')
plt.ylabel('Count')
plt.show()
```

<Figure size 864x432 with 0 Axes>




```
[43]: # Gender vs. Ratings:
plt.figure(figsize=(10, 6))
ratings_by_gender = master_data.groupby(['Gender', 'Rating']).size().unstack()
ratings_by_gender.plot(kind='bar', stacked=False, legend=True)
plt.title('Distribution of Ratings across Genders')
plt.xlabel('Gender')
plt.ylabel('Count')
plt.show()
```

<Figure size 720x432 with 0 Axes>



```
[48]: # Occupation vs. Ratings:
plt.figure(figsize=(12, 10))
ratings_by_occupation = master_data.groupby(['Occupation', 'Rating']).size().
    ↪unstack()
ratings_by_occupation.plot(kind='bar', stacked=False, legend=True)
plt.title('Distribution of Ratings across Occupations')
plt.xlabel('Occupation')
plt.ylabel('Count')
plt.show()
```

<Figure size 864x720 with 0 Axes>




```
[7]: # Calculate the number of ratings per movie :
ratings_per_movie = master_data.groupby('Title')['Rating'].count().reset_index()

# Merge with the average rating per movie
movie_ratings_stats = ratings_per_movie.merge(
    master_data.groupby('Title')['Rating'].mean().reset_index(),
    on='Title',
    how='inner',
    suffixes=('_count', '_avg')
)

# Plotting
plt.figure(figsize=(12, 6))

# Scatter plot with size representing the number of ratings
sns.scatterplot(x='Rating_avg', y='Rating_count', data=movie_ratings_stats,
               alpha=0.6, s=40)

plt.title('Number of Ratings per Movie vs. Average Rating')
plt.xlabel('Average Rating')
plt.ylabel('Number of Ratings')
plt.show()
```



```
[ ]: # 4.Develop an appropriate model to predict the movie ratings
```

```
[31]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error

# Assuming dfMaster contains your data
first_500 = master_data[:1000]

# Use the following features: 'MovieID', 'Age', 'Occupation'
features = first_500[['MovieID', 'Age', 'Occupation']].values

# Use 'Rating' as the label
labels = first_500['Rating'].values

# Split the data into training and testing sets
features_train, features_test, labels_train, labels_test = train_test_split(
    features, labels, test_size=0.2, random_state=42
)

# Decision Trees
decision_tree_model = DecisionTreeRegressor(random_state=42)
decision_tree_model.fit(features_train, labels_train)
labels_pred_decision_tree = decision_tree_model.predict(features_test)
mse_decision_tree = mean_squared_error(labels_test, labels_pred_decision_tree)
```

```

print(f'Decision Tree Mean Squared Error: {mse_decision_tree}')

# Random Forest
random_forest_model = RandomForestRegressor(n_estimators=100, random_state=42)
random_forest_model.fit(features_train, labels_train)
labels_pred_random_forest = random_forest_model.predict(features_test)
mse_random_forest = mean_squared_error(labels_test, labels_pred_random_forest)
print(f'Random Forest Mean Squared Error: {mse_random_forest}')

```

Decision Tree Mean Squared Error: 0.668740168687103  
Random Forest Mean Squared Error: 0.6553640841152201

[ ]: In both cases:  
Lower MSE values are generally better, indicating more accurate predictions on average.  
Considering the scale of your ratings, an MSE around 0.66 to 0.67 suggests that the models are making reasonably accurate predictions.

```

[14]: # Ridge Regression
from sklearn.linear_model import Ridge

# Assuming dfMaster contains your data
first_500 = master_data[:1000]

# Use the following features: 'MovieID', 'Age', 'Occupation'
features = first_500[['MovieID', 'Age', 'Occupation']].values

# Use 'Rating' as the label
labels = first_500['Rating'].values

# Split the data into training and testing sets
features_train, features_test, labels_train, labels_test = train_test_split(
    features, labels, test_size=0.2, random_state=42
)

# Ridge Regression without hyperparameter tuning
ridge_model = Ridge(alpha=1.0) # You can adjust alpha as needed
ridge_model.fit(features_train, labels_train)
labels_pred_ridge = ridge_model.predict(features_test)

# Evaluate the Ridge Regression model
mse_ridge = mean_squared_error(labels_test, labels_pred_ridge)
print(f'Ridge Regression Mean Squared Error: {mse_ridge}')

```

Ridge Regression Mean Squared Error: 0.5713180074954795



```
[ ]: Lower MSE values are generally better, indicating more accurate predictions on ↵  
↪ average.  
Ridge regression is the best model among the three models we developed here.
```